

assignment-4 September 27, 2023

0.1 Assignment - 4

0.2 RAVADAGUNDI SANATH KUMAR

sanathkumar.21bce9610@vitapstudent.ac.in
(<mailto:sanathkumar.21bce9610@vitapstudent.ac.in>)

0.3 • Data Preprocessing.

- o Import the Libraries.
- o Importing the dataset.
- o Checking for Null Values.
- o Data Visualization.
- o Outlier Detection
- o Splitting Dependent and Independent variables
- o- Encoding
- o Feature Scaling.
- o Splitting Data into Train and Test.

0.3.1 Import the Libraries.

```
In [72]: ▶ import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

0.3.2 Importing the dataset.


```
In [73]: ▶ df=pd.read_csv("C:/Users/rsana/Downloads/archive/WA_Fn-UseC_-HR-Employee-Attr
```

In [74]: `df.head()`

Out[74]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Edu
0	41	Yes	Travel_Rarely	1102	Sales	1	2	L
1	49	No	Travel_Frequently	279	Research & Development	8	1	L
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	L
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 35 columns



In [75]: `df.shape`

Out[75]: (1470, 35)

In [76]: `df.info()`


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    1470 non-null   int64
1   Attrition                            1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                            1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                    1470 non-null   int64
6   Education                            1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                       1470 non-null   int64
9   EmployeeNumber                      1470 non-null   int64
10  EnvironmentSatisfaction              1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                      1470 non-null   int64
14  JobLevel                            1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                     1470 non-null   int64
17  MaritalStatus                       1470 non-null   object
18  MonthlyIncome                       1470 non-null   int64
19  MonthlyRate                         1470 non-null   int64
20  NumCompaniesWorked                  1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                   1470 non-null   int64
24  PerformanceRating                   1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                       1470 non-null   int64
27  StockOptionLevel                    1470 non-null   int64
28  TotalWorkingYears                   1470 non-null   int64
29  TrainingTimesLastYear               1470 non-null   int64
30  WorkLifeBalance                     1470 non-null   int64
31  YearsAtCompany                      1470 non-null   int64
32  YearsInCurrentRole                  1470 non-null   int64
33  YearsSinceLastPromotion              1470 non-null   int64
34  YearsWithCurrManager                 1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

```
In [77]: df.describe()
```

Out[77]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeN
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.0
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.8
std	9.135373	403.509100	8.106864	1.024165	0.0	602.0
min	18.000000	102.000000	1.000000	1.000000	1.0	1.0
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.2
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.8
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.7
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.0

8 rows × 26 columns



0.3.3 Checking for null values

```
In [78]: ▶ df.isnull().any()
```

```
Out[78]: Age                False
Attrition                False
BusinessTravel           False
DailyRate               False
Department              False
DistanceFromHome        False
Education               False
EducationField           False
EmployeeCount            False
EmployeeNumber           False
EnvironmentSatisfaction  False
Gender                  False
HourlyRate              False
JobInvolvement           False
JobLevel                False
JobRole                 False
JobSatisfaction          False
MaritalStatus            False
MonthlyIncome            False
MonthlyRate             False
NumCompaniesWorked       False
Over18                  False
OverTime                 False
PercentSalaryHike        False
PerformanceRating        False
RelationshipSatisfaction  False
StandardHours            False
StockOptionLevel         False
TotalWorkingYears        False
TrainingTimesLastYear    False
WorkLifeBalance          False
YearsAtCompany           False
YearsInCurrentRole       False
YearsSinceLastPromotion  False
YearsWithCurrManager     False
dtype: bool
```

```
In [79]: ► df.isnull().sum()
```

```
Out[79]: Age                                0
Attrition                                0
BusinessTravel                           0
DailyRate                                0
Department                               0
DistanceFromHome                         0
Education                                0
EducationField                           0
EmployeeCount                            0
EmployeeNumber                           0
EnvironmentSatisfaction                  0
Gender                                    0
HourlyRate                               0
JobInvolvement                           0
JobLevel                                 0
JobRole                                  0
JobSatisfaction                          0
MaritalStatus                            0
MonthlyIncome                            0
MonthlyRate                              0
NumCompaniesWorked                       0
Over18                                    0
OverTime                                  0
PercentSalaryHike                        0
PerformanceRating                        0
RelationshipSatisfaction                  0
StandardHours                            0
StockOptionLevel                         0
TotalWorkingYears                        0
TrainingTimesLastYear                    0
WorkLifeBalance                          0
YearsAtCompany                           0
YearsInCurrentRole                       0
YearsSinceLastPromotion                   0
YearsWithCurrManager                     0
dtype: int64
```

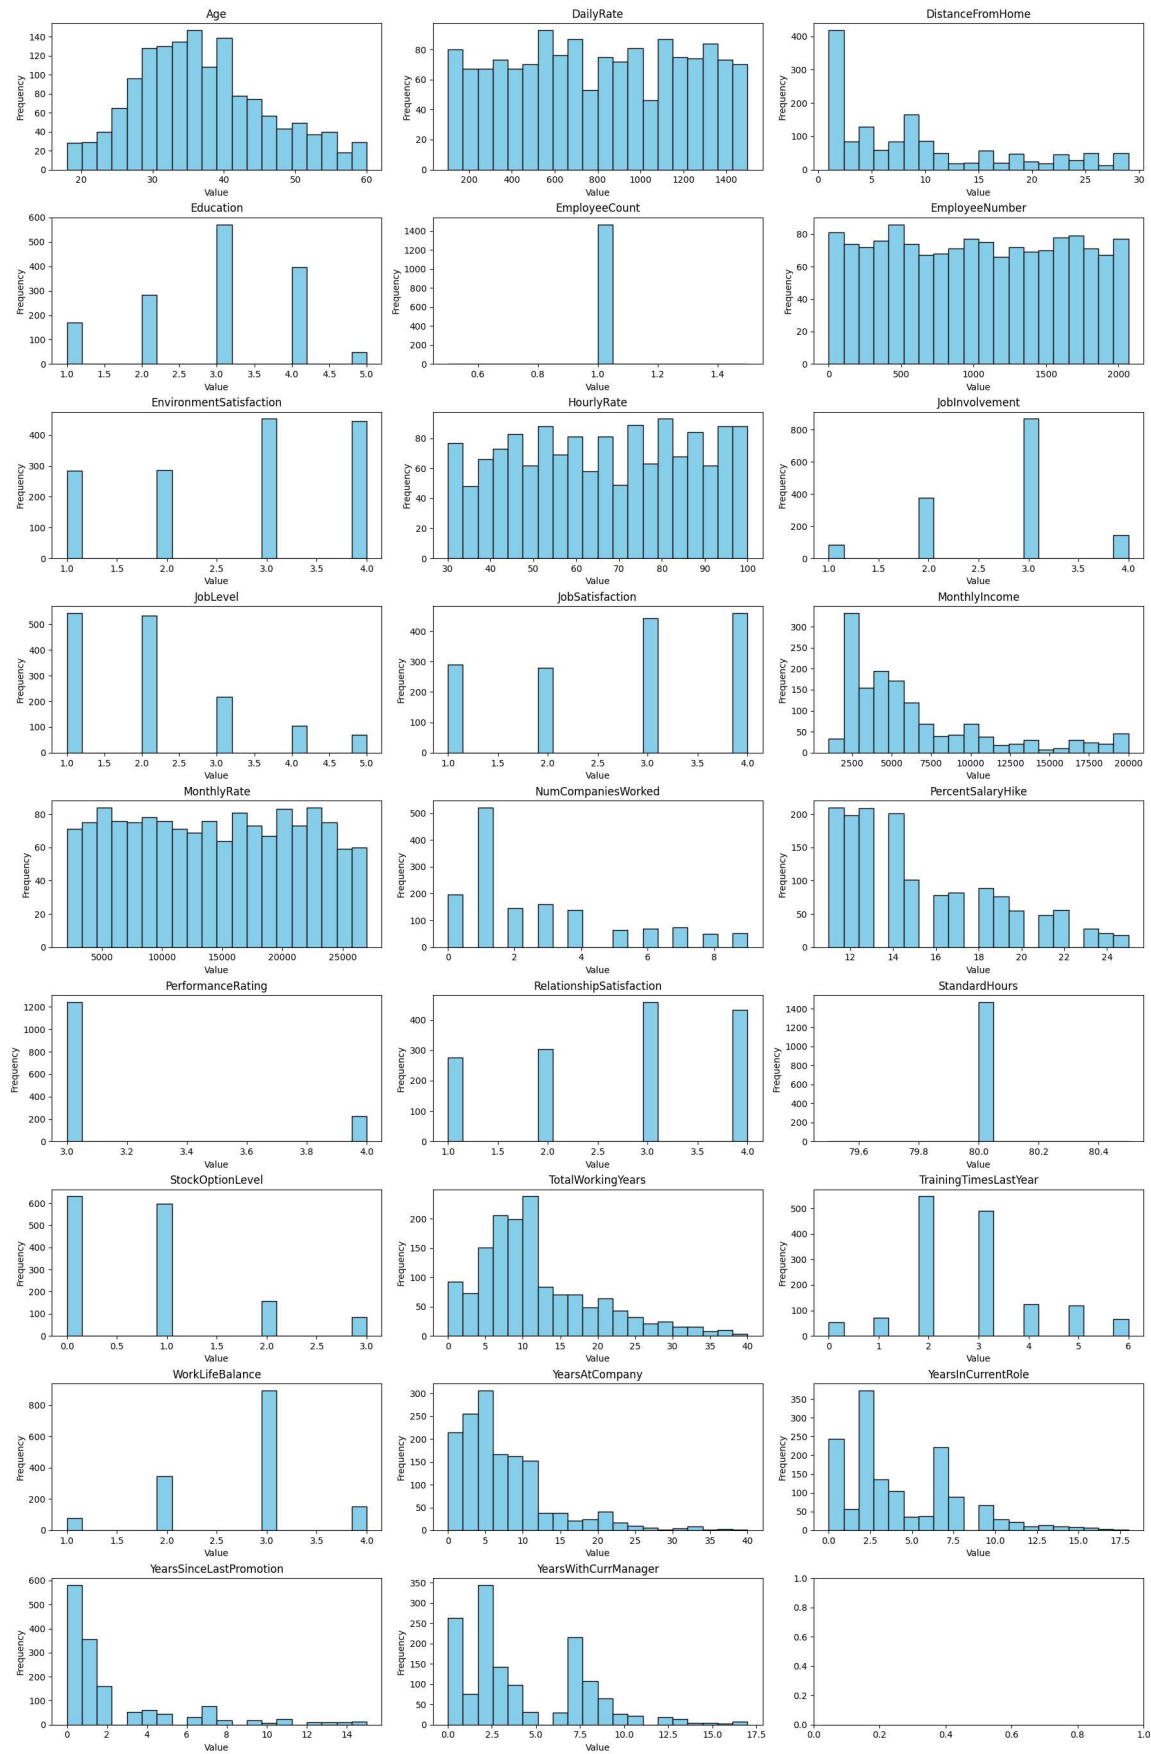
```
In [80]: ► import warnings
# Ignore the specified warnings globally
warnings.filterwarnings("ignore", category=FutureWarning)
warnings.filterwarnings("ignore", category=UserWarning)
```

0.3.4 Data Visualization

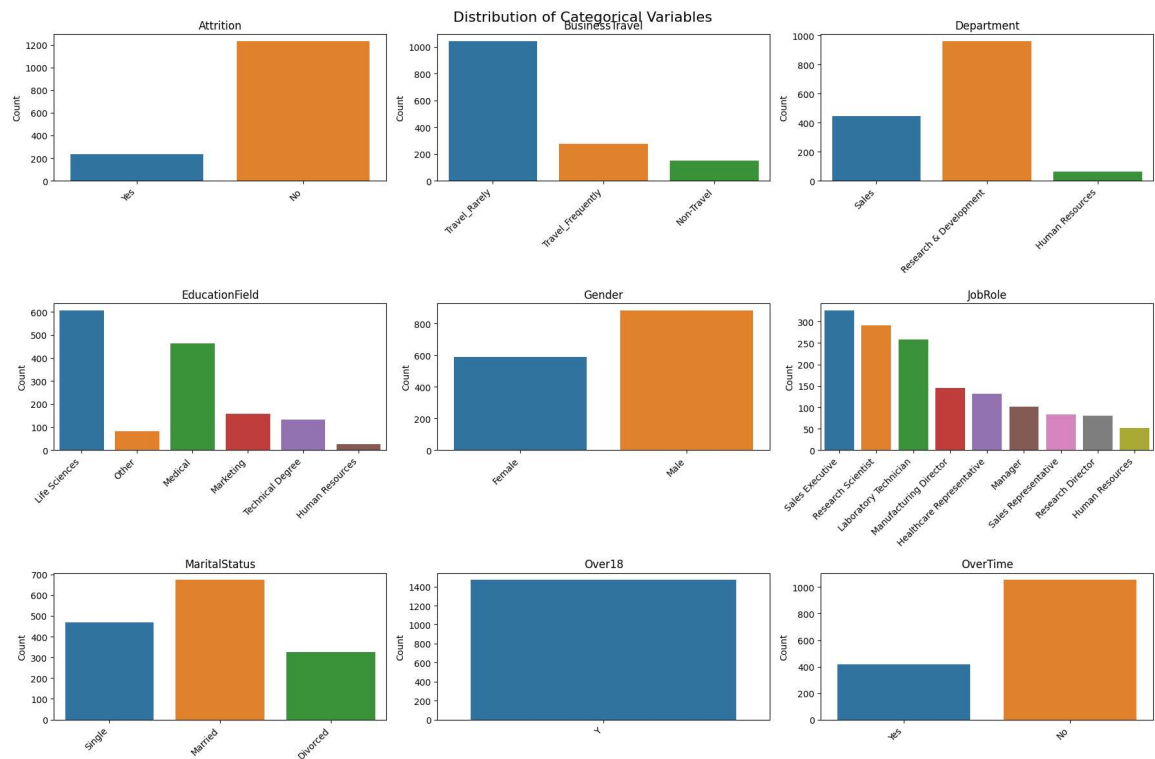
```
In [81]: ▶ numerical_vars = df.select_dtypes(include=['int64', 'float64']).columns.tolist
# Create histograms for the selected numerical variables
fig, axes = plt.subplots(nrows=9, ncols=3, figsize=(18, 30))
fig.suptitle('Distribution of Numerical Variables', fontsize=16)
for ax, var in zip(axes.flatten(), numerical_vars):
    ax.hist(df[var], bins=20, color='skyblue', edgecolor='black')
    ax.set_title(var)
    ax.set_xlabel('Value')
    ax.set_ylabel('Frequency')

plt.tight_layout()
plt.subplots_adjust(top=0.9)
plt.show()
```

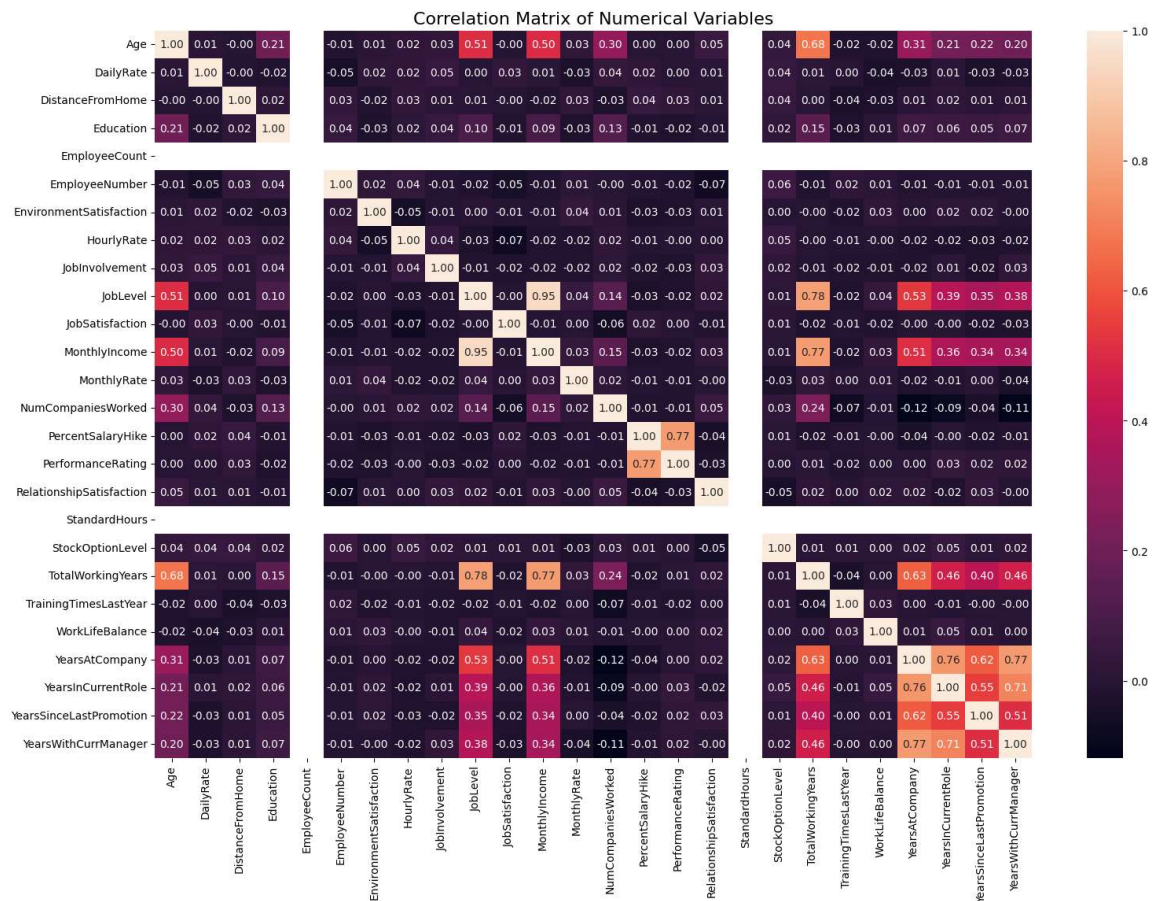
Distribution of Numerical Variables




```
In [82]: # Select categorical variables for visualization
categorical_vars = df.select_dtypes(include=['object']).columns.tolist()
# Create bar plots for the selected categorical variables
fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(18, 12))
fig.suptitle('Distribution of Categorical Variables', fontsize=16)
for ax, var in zip(axes.flatten(), categorical_vars):
    sns.countplot(data=df, x=var, ax=ax)
    ax.set_title(var)
    ax.set_xlabel('')
    ax.set_ylabel('Count')
    ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')
# Remove the empty subplots
for ax in axes.flatten()[len(categorical_vars):]:
    fig.delaxes(ax)
plt.tight_layout()
plt.subplots_adjust(top=0.95)
plt.show()
```



```
In [83]: # Compute the correlation matrix
correlation_matrix = df.corr(numeric_only=True)
# Plot the heatmap of the correlation matrix
plt.figure(figsize=(18, 12))
sns.heatmap(correlation_matrix, annot=True, fmt='.2f')
plt.title('Correlation Matrix of Numerical Variables', fontsize=16)
plt.show()
```



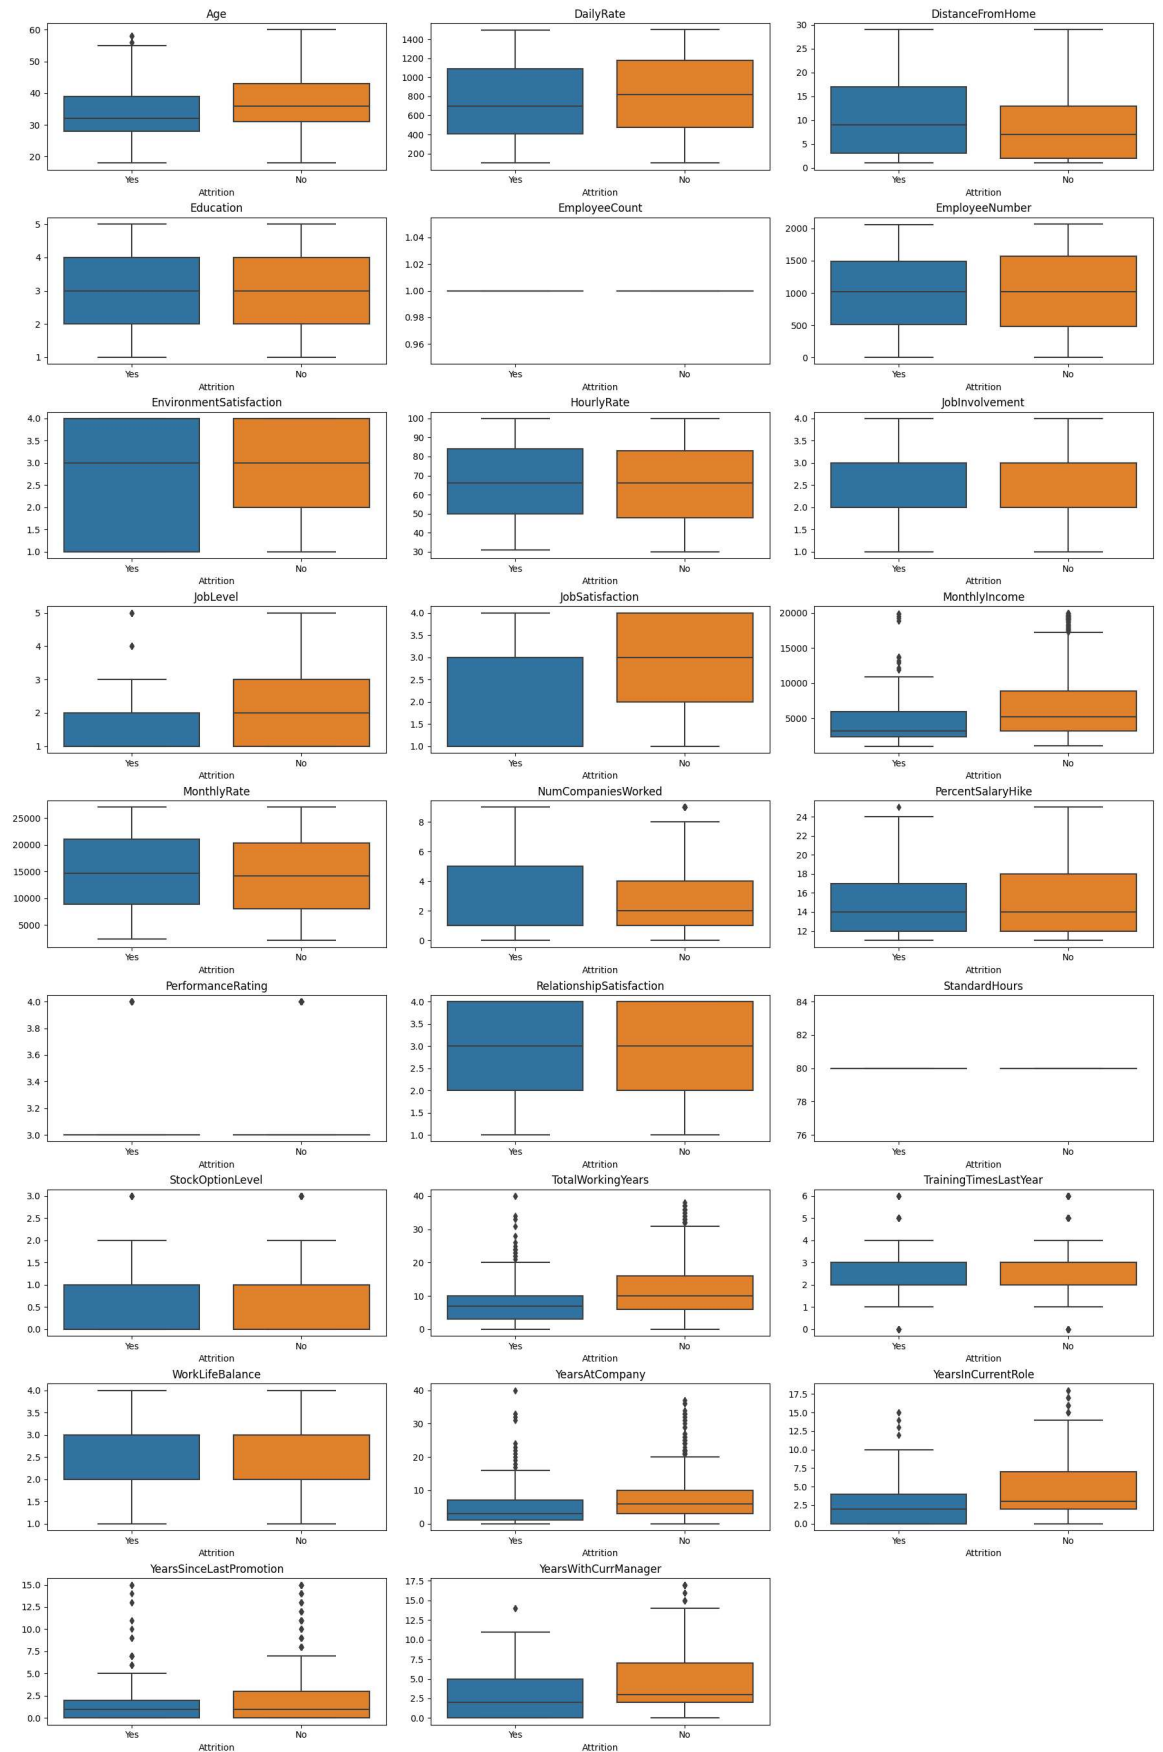
```
In [84]: ▶ # Select numerical variables for visualization
numerical_vars_to_compare = df.select_dtypes(include=['int64', 'float64']).columns

# Create box plots for the selected numerical variables grouped by Attrition
fig, axes = plt.subplots(nrows=9, ncols=3, figsize=(18, 30))
fig.suptitle('Distribution of Numerical Variables by Attrition', fontsize=16)

for ax, var in zip(axes.flatten(), numerical_vars_to_compare):
    sns.boxplot(data=df, x='Attrition', y=var, ax=ax)
    ax.set_title(var)
    ax.set_xlabel('Attrition')
    ax.set_ylabel('')

# Remove the empty subplot
fig.delaxes(axes.flatten()[-1])
plt.tight_layout()
plt.subplots_adjust(top=0.9)
plt.show()
```

Distribution of Numerical Variables by Attrition



0.3.5 Splitting Dependent and Independent variables

```
In [85]: x = df.drop(columns=['Attrition'],axis=1)
         y = df['Attrition']
```

```
In [86]: x.shape,y.shape
```

```
Out[86]: ((1470, 34), (1470,))
```

```
In [87]: y.head()
```

```
Out[87]: 0    Yes
         1    No
         2    Yes
         3    No
         4    No
         Name: Attrition, dtype: object
```

0.3.6 Encoding

```
In [88]: from sklearn.preprocessing import LabelEncoder
         le=LabelEncoder()
```

```
In [89]: y=le.fit_transform(y)
```

```
In [90]: x["BusinessTravel"]=le.fit_transform(x["BusinessTravel"])
         x["Department"]=le.fit_transform(x["Department"])
         x["EducationField"]=le.fit_transform(x["EducationField"])
         x["Gender"]=le.fit_transform(x["Gender"])
         x["JobRole"]=le.fit_transform(x["JobRole"])
         x["MaritalStatus"]=le.fit_transform(x["MaritalStatus"])
         x["Over18"]=le.fit_transform(x["Over18"])
         x["OverTime"]=le.fit_transform(x["OverTime"])
```

In [91]: `x.head()`

Out[91]:

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField
0	41	2	1102	2	1	2	1
1	49	1	279	1	8	1	1
2	37	2	1373	1	2	2	4
3	33	1	1392	1	3	4	1
4	27	2	591	1	2	1	3

5 rows × 34 columns

0.3.7 Feature Scaling

In [92]: `from sklearn.preprocessing import MinMaxScaler
ms = MinMaxScaler()`

In [93]: `x_scaled = pd.DataFrame(ms.fit_transform(x), columns = x.columns)
x_scaled.head()`

Out[93]:

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationF
0	0.547619	1.0	0.715820	1.0	0.000000	0.25	
1	0.738095	0.5	0.126700	0.5	0.250000	0.00	
2	0.452381	1.0	0.909807	0.5	0.035714	0.25	
3	0.357143	0.5	0.923407	0.5	0.071429	0.75	
4	0.214286	1.0	0.350036	0.5	0.035714	0.00	

5 rows × 34 columns

0.3.8 Splitting into Training and Testing dataset

In [94]: `from sklearn.model_selection import train_test_split`

In [95]: `x_train1, x_test1, y_train1, y_test1 = train_test_split(x, y,
test_size = 0.3, random_state = 0)`

```
In [96]: ▶ print(x_train1.shape)
          print(x_test1.shape)
          print(y_train1.shape)
          print(y_test1.shape)
```

```
(1029, 34)
(441, 34)
(1029,)
(441,)
```

```
In [97]: ▶ x_train2, x_test2, y_train2, y_test2 = train_test_split(x_scaled, y,
          test_size = 0.3, random_state = 0)
```

```
In [98]: ▶ print(x_train2.shape)
          print(x_test2.shape)
          print(y_train2.shape)
          print(y_test2.shape)
```

```
(1029, 34)
(441, 34)
(1029,)
(441,)
```

0.4 Model Building

- 1.Import the model building Libraries
- 2.Initializing the model
- 3.Training and testing the model
- 4.Evaluation of Model
- 5.Save the Model

0.4.1 Logistic Regression

```
In [99]: ▶ from sklearn.linear_model import LogisticRegression
          lrmodel = LogisticRegression()
```

```
In [100]: ▶ lrmodel.fit(x_train2,y_train2)
```

```
Out[100]: ▾ LogisticRegression
           LogisticRegression()
```

```
In [101]: ▶ predlr=lrmodel.predict(x_test2)
```

In [102]: `from sklearn.metrics import accuracy_score, confusion_matrix, classification_r`

In [103]: `accuracy_score(y_test2, predlr)`

Out[103]: 0.8866213151927438

In [104]: `confusion_matrix(y_test2, predlr)`

Out[104]: array([[368, 3],
[47, 23]], dtype=int64)

In [105]: `pd.crosstab(y_test2, predlr)`

Out[105]:

	col_0	0	1
row_0			
0	368	3	
1	47	23	

In [106]: `print(classification_report(y_test2, predlr))`


	precision	recall	f1-score	support
0	0.89	0.99	0.94	371
1	0.88	0.33	0.48	70
accuracy			0.89	441
macro avg	0.89	0.66	0.71	441
weighted avg	0.89	0.89	0.86	441

In [107]: `roc_auc_score(y_test2, predlr)`

Out[107]: 0.6602425876010781

In [108]: `prob1 = lrmodel.predict_proba(x_test2)[: , 1]
fpr1, tpr1, thres1 = roc_curve(y_test2, prob1)`


```
In [109]: ▶ plt.plot(fpr1, tpr1)
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title("ROC Curve")
plt.show()
```



0.4.2 Decision Tree

```
In [110]: ▶ from sklearn.tree import DecisionTreeClassifier
dtmodel=DecisionTreeClassifier()
```

```
In [111]: ▶ dtmodel.fit(x_train1, y_train1)
```

```
Out[111]: ▼ DecisionTreeClassifier
DecisionTreeClassifier()
```

```
In [112]: ▶ prededt = dtmodel.predict(x_test1)
```

```
In [113]: ▶ accuracy_score(y_test1, prededt)
```

```
Out[113]: 0.764172335600907
```

```
In [114]: ► confusion_matrix(y_test1,preddt)
```

```
Out[114]: array([[315,  56],
                 [ 48,  22]], dtype=int64)
```

```
In [115]: ► pd.crosstab(y_test1,preddt)
```

```
Out[115]:
```

col_0	0	1
row_0		
0	315	56
1	48	22

```
In [116]: ► print(classification_report(y_test1,preddt))
```


	precision	recall	f1-score	support
0	0.87	0.85	0.86	371
1	0.28	0.31	0.30	70
accuracy			0.76	441
macro avg	0.57	0.58	0.58	441
weighted avg	0.77	0.76	0.77	441

```
In [117]: ► roc_auc_score(y_test1, preddt)
```

```
Out[117]: 0.5816711590296496
```

```
In [118]: ► prob2 = dtmodel.predict_proba(x_test1)[:,-1]
fpr2, tpr2, thres2 = roc_curve( y_test1, prob2)
```

```
In [119]: ▶ plt.plot(fpr2, tpr2)
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title("ROC Curve")
plt.show()
```



0.4.3 Random Forest

```
In [120]: ▶ from sklearn.ensemble import RandomForestClassifier
rfmodel = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
In [121]: ▶ rfmodel.fit(x_train1, y_train1)
```

```
Out[121]: ▼      RandomForestClassifier
RandomForestClassifier(random_state=42)
```

```
In [122]: ▶ predrf=rfmodel.predict(x_test1)
accuracy_score(y_test1,predrf)
```

```
Out[122]: 0.8616780045351474
```

In [123]: `confusion_matrix(y_test1, predrf)`

Out[123]: `array([[367, 4],
 [57, 13]], dtype=int64)`

In [124]: `pd.crosstab(y_test1, predrf)`

Out[124]:

col_0	0	1
row_0		
0	367	4
1	57	13

In [125]: `print(classification_report(y_test1, predrf))`

	precision	recall	f1-score	support
0	0.87	0.99	0.92	371
1	0.76	0.19	0.30	70
accuracy			0.86	441
macro avg	0.82	0.59	0.61	441
weighted avg	0.85	0.86	0.82	441

0.4.4 Accuracy of All Three Models

In [126]: `print("Accuracy of Logistic Regression :", accuracy_score(y_test2, predlr))
print("Accuracy of Decision Tree Classifier :", accuracy_score(y_test1, predt))
print("Accuracy of Random Forest Classifier :", accuracy_score(y_test1, predrf))`

Accuracy of Logistic Regression : 0.8866213151927438
Accuracy of Decision Tree Classifier : 0.764172335600907
Accuracy of Random Forest Classifier : 0.8616780045351474